

# A gentle introduction to the art of exploitation

Part 1 : Buffer overflow 101

by anticlockwise & malweisse

# About us

**Andrea Fioraldi** [ @andrea Fioraldi ]

**Pietro Borrello** [ @pietroborrello ]

Cyberchallengers 2017

We founded TheRomanXploit and started  
pwning right after

Then mHACKeroni with Milan, Venice and  
Padua guys, to participate to DEFCON CTF

← Microarchitectural attacks

Fuzzing and binary analysis →



# What can possibly go wrong?

what do you see?

0x58 0x33 0x58 0x30

# What can possibly go wrong?

what do you see?

0x58 0x33 0x58 0x30



ascii?

```
>>> '58335830'.decode('hex')  
'X3X0'
```

# What can possibly go wrong?

what do you see?

0x58 0x33 0x58 0x30



an integer?

0x30583358 = 811086680

# What can possibly go wrong?

what do you see?

0x58 0x33 0x58 0x30



an address?

0x30583358 <main>

# What can possibly go wrong?

what do you see?

0x58 0x33 0x58 0x30



code?

0: 58

pop rax

1: 33 58 30

xor ebx,DWORD PTR [rax+0x30]

# What can possibly go wrong?

what do you see?

0x58 0x33 0x58 0x30



kim?





# What can possibly go wrong?

what do you see?

0x58 0x33 0x58 0x30



they are just bytes!

0x58 0x33 0x58 0x30

It is just the context that gives the meaning to the bytes

# What can possibly go wrong?

0x58 0x33 0x58 0x30

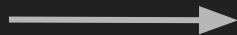
Any time you mess with the context you open the window  
to exploitation

A vulnerability allows us to make the CPU manage bytes in  
the wrong context

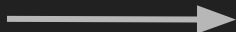
# Buffer overflow

```
void foo() {  
    char buffer[10];  
    scanf("%s", buffer); //DOH  
}
```

DATA  
context



ADDRESS  
context



buffer (10 bytes)

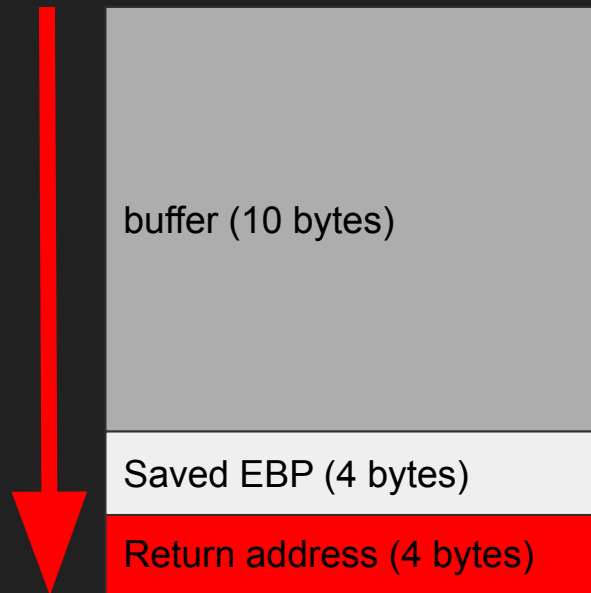
Saved EBP (4 bytes)

Return address (4 bytes)

# Buffer overflow

Let's override the return  
address content!

After foo() the execution will be  
resumed from the return  
address.



# Buffer overflow

For example:

“aaaaaaaaaa” + “aaaa” + addr\_to\_return

↑                    ↑  
padding (10)    ebp

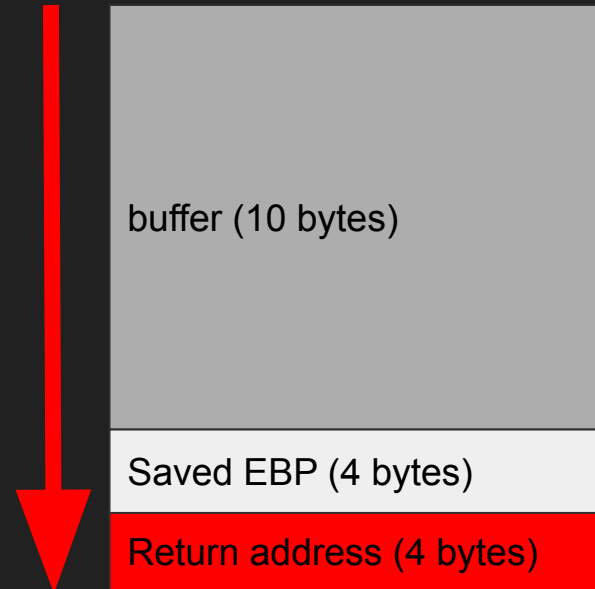


# Buffer overflow

```
void foo() {  
  
    char buffer[10];  
  
    scanf("%s", buffer); //DOH  
  
}
```

⇒ return anywhere you want!

but... Where?



# Data or Code? Shellcode!

- Mix up code and data is the worst possible vulnerability!
- A shellcode is the payload you manage to execute
  1. Write the shellcode in memory (as DATA context)
    - es: read(0, buffer, 0x1000)
  2. Redirect the execution to the shellcode (as CODE context)
    - es: overwrite the return address with a vulnerability
  3. WIN
- Assuming you can write in a RWX page... (Who said “please switch to ROP”?)

# Shellcode?

- Assume:
  1. You can write to RWX memory
  2. You can override the return address to jump to RWX memory
- What do you write?



# Shellcode?

- Assume:
  1. You can write to RWX memory
  2. You can override the return address to jump to RWX memory
- What do you write?

⇒ Usually the goal is to spawn a shell in the victim server

`system( "/bin/sh" )` will do the job!

But, how do you call it, if it is not present in the binary?

⇒ Implement it with your shellcode!

# Spawn a Shell!

Let's dissect system!

Browse the source code of [glibc/sysdeps/posix/system.c](#)  
do\_system

```
116 #ifdef FORK
117     pid = FORK ();
118 #else
119     pid = __fork ();
120 #endif
121     if (pid == (pid_t) 0)
122     {
123         /* Child side. */
124         const char *new_argv[4];
125         new_argv[0] = SHELL_NAME;
126         new_argv[1] = "-c";
127         new_argv[2] = line;
128         new_argv[3] = NULL;
129
130         /* Restore the signals. */
131         (void) __sigaction (SIGINT, &intr, (struct sigaction *) NULL);
132         (void) __sigaction (SIGQUIT, &quit, (struct sigaction *) NULL);
133         (void) __sigprocmask (SIG_SETMASK, &omask, (sigset_t *) NULL);
134         INIT_LOCK ();
135
136         /* Exec the shell. */
137         (void) __execve (SHELL_PATH, (char *const *) new_argv, __environ);
138         _exit (127);
139     }
```

# Spawn a Shell!

`__execve` is the stub for the `execve` system call: the system call that transforms the calling process into a new process to execute.

It sets for **32 bits**:

`eax` -> **0x0b**  
`ebx` -> address of `"/bin/sh"`  
`ecx` -> NULL  
`edx` -> NULL

then executes:

`int` **0x80**

or for **64 bits**:

`rax` -> **0x3b**  
`rdi` -> address of `"/bin/sh"`  
`rsi` -> NULL  
`rdx` -> NULL

then executes:

`syscall`

# Spawn a Shell!

Therefore our shellcode will be something like:

```
push 0x68732f    ←  
push 0x6e69622f ← Write /bin/sh into the stack  
mov     ebx, esp  
mov     ecx, 0x0  
mov     edx, 0x0  
mov     eax, 0xb  
int     0x80
```

# Spawn a Shell!

Write the shellcode in RWX memory and jump to it!

```
push 0x68732f
push 0x6e69622f
mov   ebx, esp
mov   ecx, 0x0
mov   edx, 0x0
mov   eax, 0xb
int   0x80
```

buffer (10 bytes)

Saved EBP (4 bytes)

Return address (4 bytes)

# Are shellcodes dead?

- All modern OS enforce NX protection:
  - No region in the binary can be writable and executable at the same time when compiling -> no RWX regions!

# Are shellcodes dead?

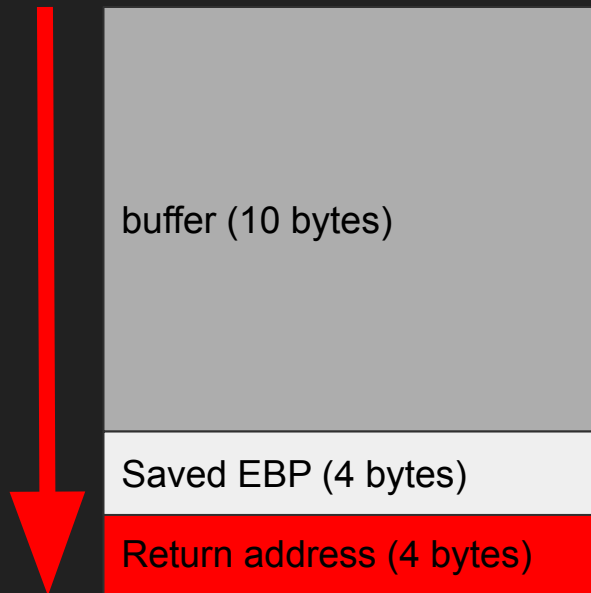
- All modern OS enforce NX protection:
  - No region in the binary can be writable and executable at the same time when compiling -> no RWX regions!
- But any self modifying program will need writable and executable pages!
  - Browser JIT compilers
  - Browser WASM compilers
  - QEMU (sometimes not enforces neither ASLR nor NX)

⇒ Still useful to know how shellcodes work

# Buffer overflow returned

```
void foo() {  
  
    char buffer[10];  
  
    scanf("%s", buffer); //DOH  
  
}
```

⇒ Now we have NX so we cannot jump to shellcode!



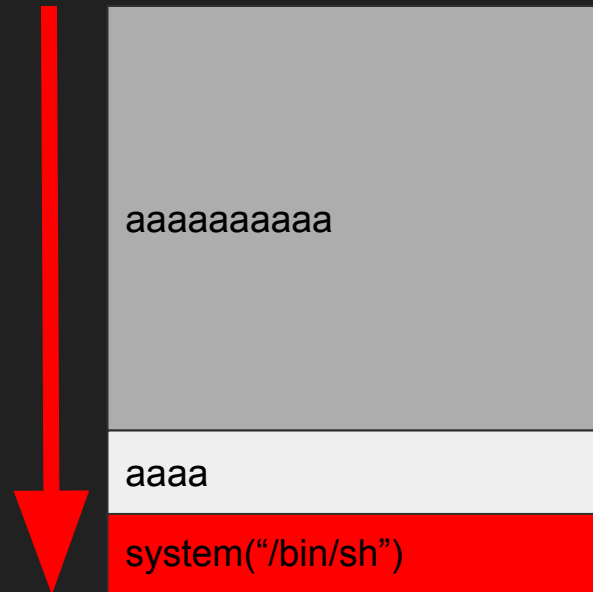


# Buffer overflow

We can still put anything we want in  
`addr_to_return!`

Calling `system("/bin/sh")` will execute a shell  
in the current context

But what if we don't have `system` available to  
call?



# Spawn a Shell!

We need to execute something similar to:

```
mov    ebx, /bin/sh_string_address
mov    ecx, 0x0
mov    edx, 0x0
mov    eax, 0xb
int    0x80
```

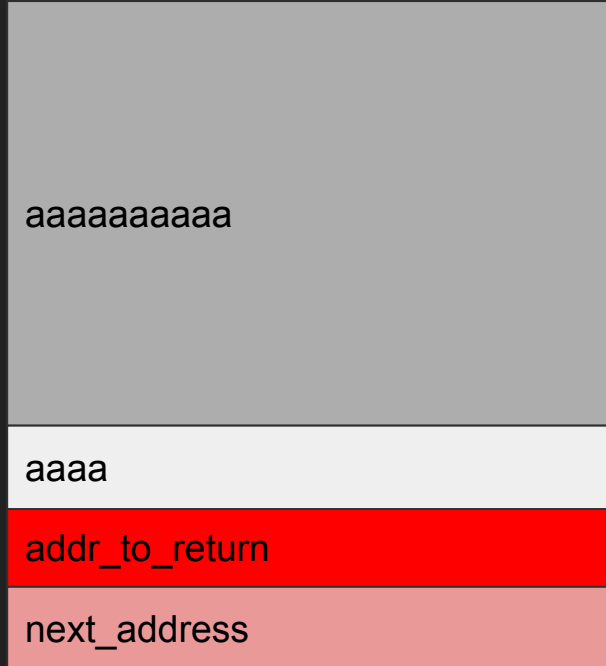
But how?

... Let's notice something

# Composing a chain

If `addr_to_return` point to some code that ends in “ret”

the execution will then continue from the address after



...and we can iterate again and again

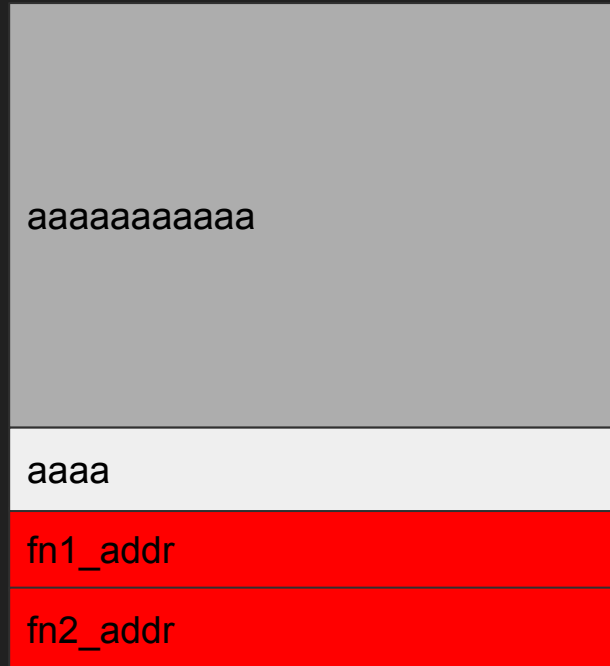
# Composing a chain

So why we don't insert more than one address to return to?

For example:

“aaaaaaaaaa” + “aaaa” + fn1\_addr + fn2\_addr

↑            ↑  
padding (10)    ebp



# Composing a chain

fn1:

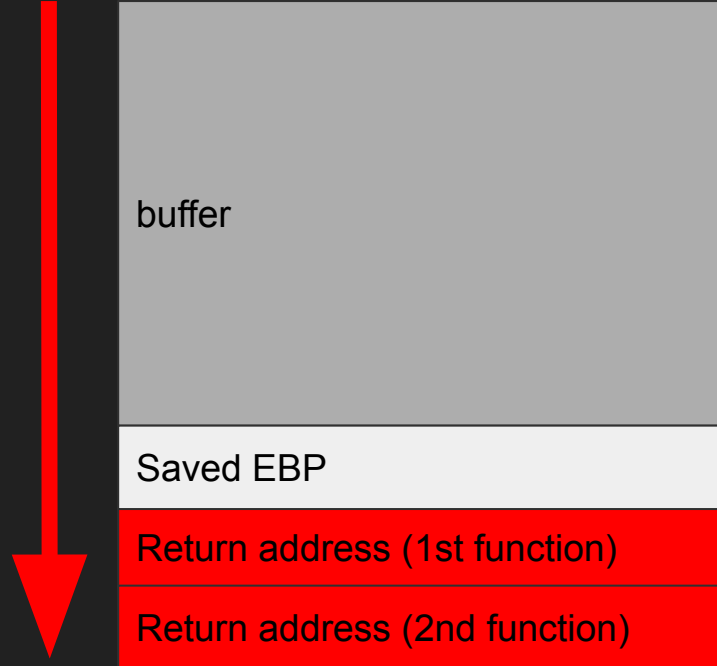
```
mov eax, 0xabadcafe
```

```
ret
```

fn2:

```
mov ecx, eax
```

```
ret
```



# Composing a chain

foo:

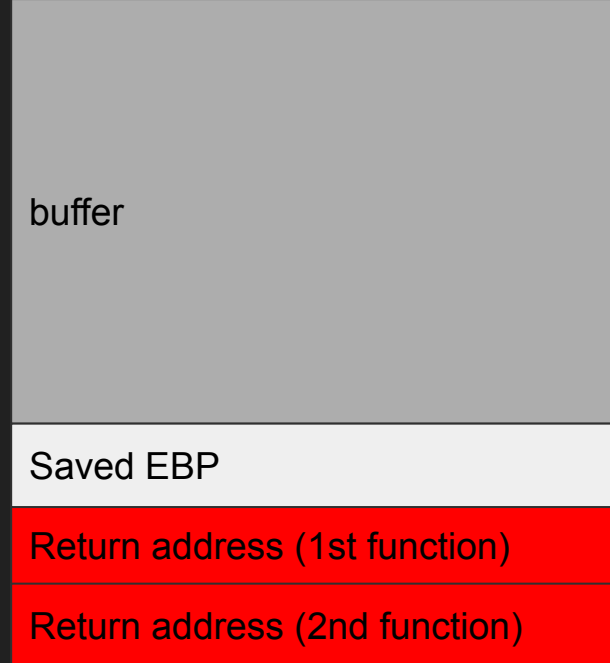
...

ret ← EIP

EAX: ?

ECX: ?

ESP →



# Composing a chain

fn1:

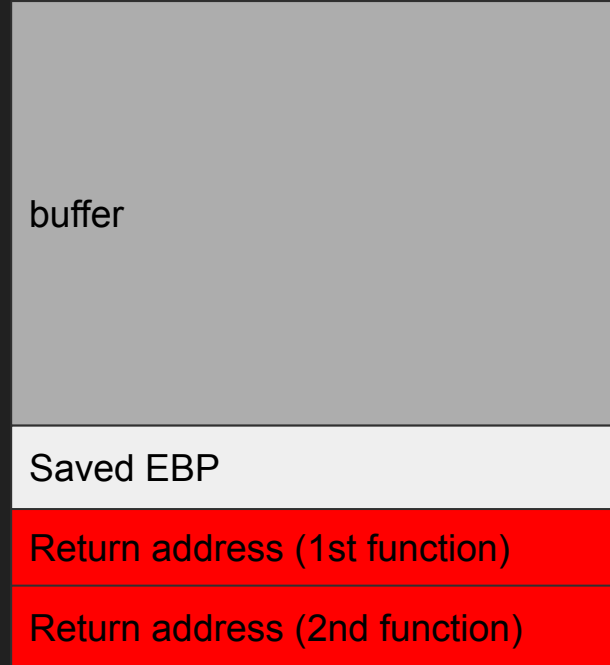
mov eax, 0xabadcafe ← EIP

ret

EAX: ?

ECX: ?

ESP →



# Composing a chain

fn1:

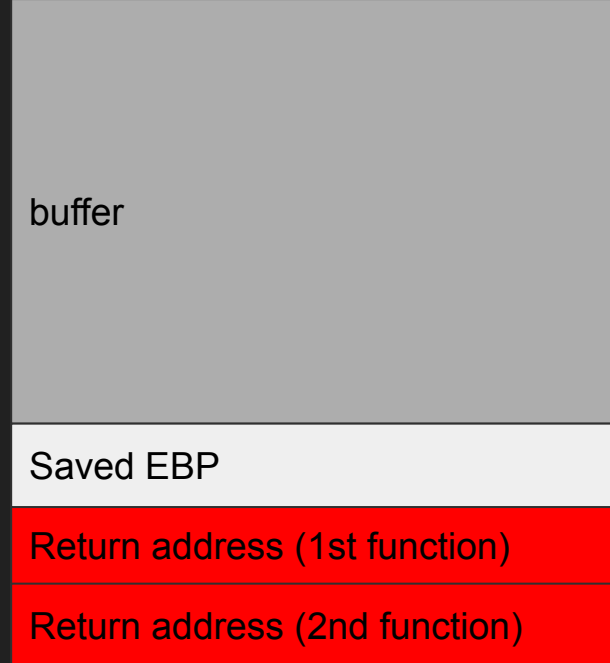
mov eax, 0xabadcafe

ret ← EIP

EAX: 0xabadcafe

ECX: ?

ESP →





# Composing a chain

fn2:

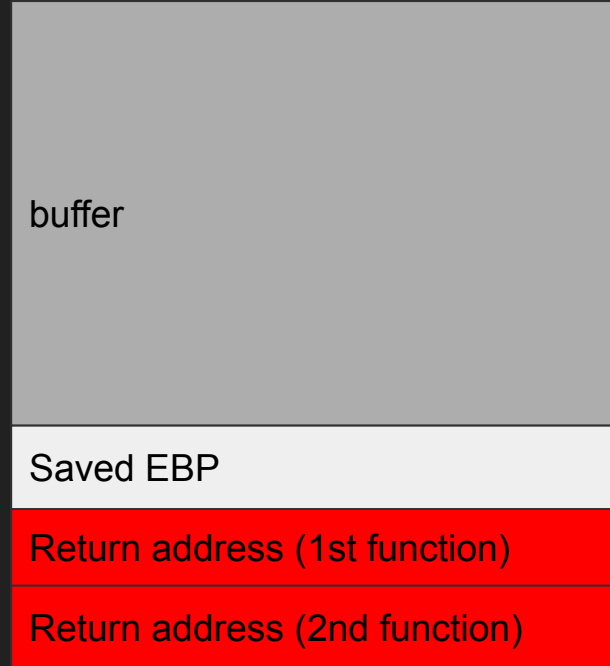
mov ecx, eax ← EIP

ret

EAX: 0xabadcafe

ECX: ?

ESP →



# Composing a chain

fn2:

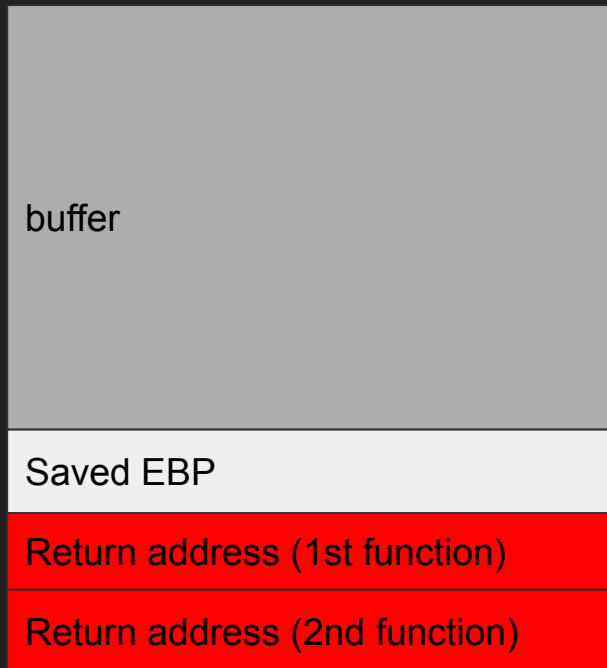
```
mov ecx, eax
```

```
ret ← EIP
```

EAX: 0xabadcafe

ECX: 0xabadcafe

ESP →



# Composing a chain

We can insert addresses to return to execute, as long as we want!

BUT: every piece of code we execute must end with a ret instruction, to continue the chain

Return address (1st function)

Return address (2nd function)

Return address (3rd function)

Return address (4th function)

Return address (5th function)

# Composing a chain

For example, let's load ebx with the address of a /bin/sh string

```
mov ebx, writable_address; ret;
```

```
pop eax; ret;
```

```
"/bin"
```

```
mov [ebx], eax; ret;
```

```
pop eax; ret;
```

```
"/sh"
```

```
mov [ebx+4], eax; ret;
```

Now ebx points to the /bin/sh string

x86 instructions are not aligned:

mov ah, 90; ret ---> **b4 5a c3**

pop edx; ret ---> **5a c3**

Every sequence of bytes endings with ret (**0xc3**) can be used to build a chain.

Find these sequences with tools like ROPGadget or ropper.

```
$ ROPGadget --binary <bin_to_pwn>
```

or

```
$ ropper -f <bin_to_pwn>
```

# Existing Countermeasures

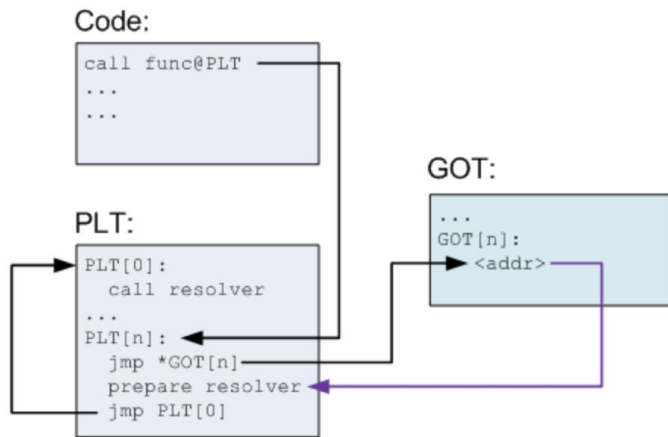
- ASLR
- Stack Canary
- NX
- PIE

## Address Space Layout Randomization

is a protection technique that randomizes  
the base address of stack, heap and  
library code



# The GOT and the PLT



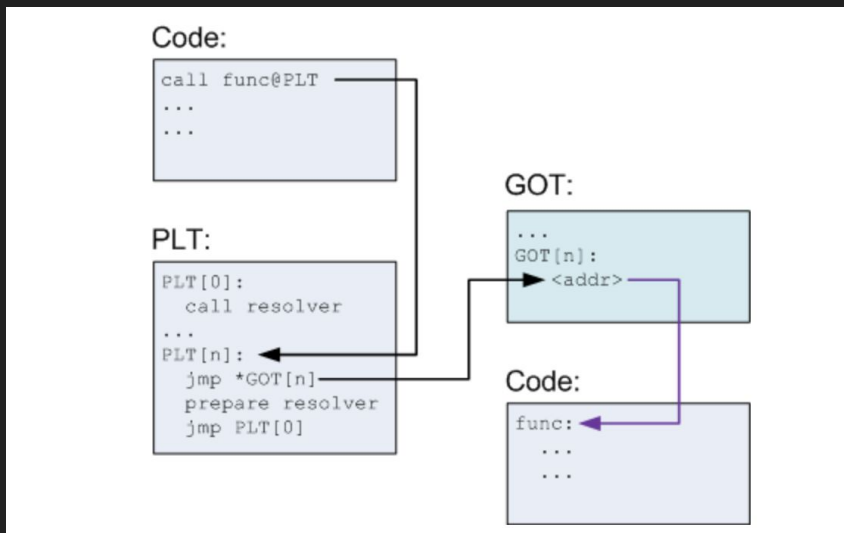
when called for the first time, `func@PLT` will jump to the corresponding

`got_entry:<func_loader_stub>`

that will call the loader for the function



# The GOT and the PLT



then the loader will replace the entry with the right address, consulting at runtime the symbol map of the library, to resolve the randomized address

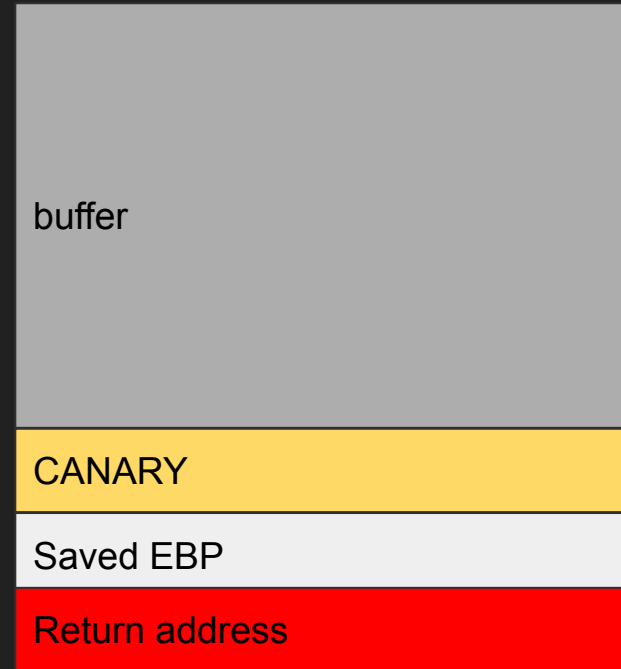
got\_entry: <address of func in libc>

# Stack Canary

A random value is automatically inserted before any return address and checked that remained the same before doing any ret instruction.

To modify the return address need to overwrite the canary.

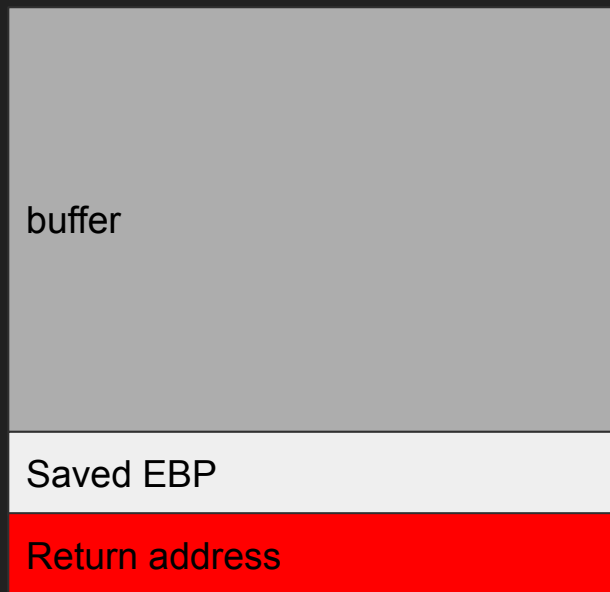
The program is crashed if the canary is modified!



No region in the binary can be writable and executable at the same time when compiling.

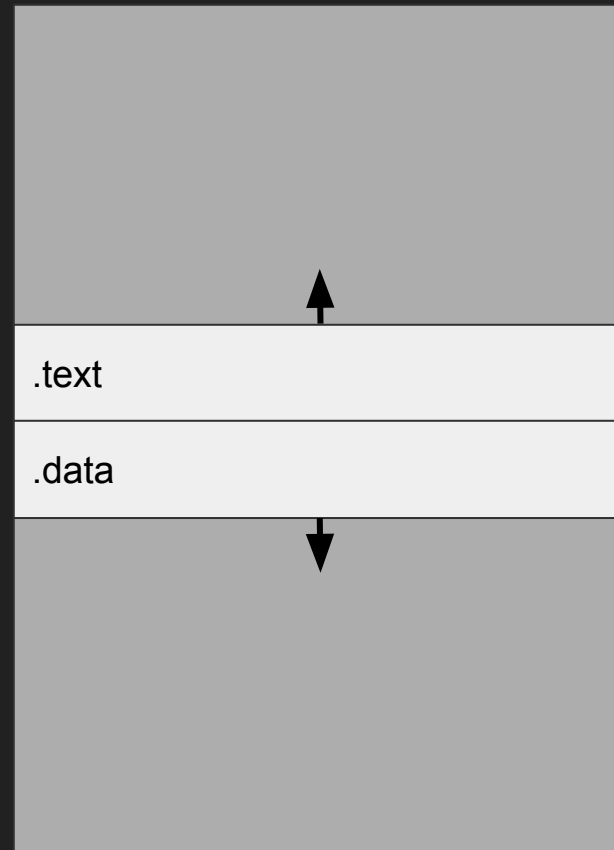
ROP is not affected by this protection in any way ;)

Only against classical shellcodes



Position Independent Executables (PIE) are binaries made entirely by position independent code. The base of the whole executable is loaded at random memory position.

ROP mitigation, since need to discover the base address before the exploit can begin.



# Other Protections

- Partial RelRO
- Full RelRO

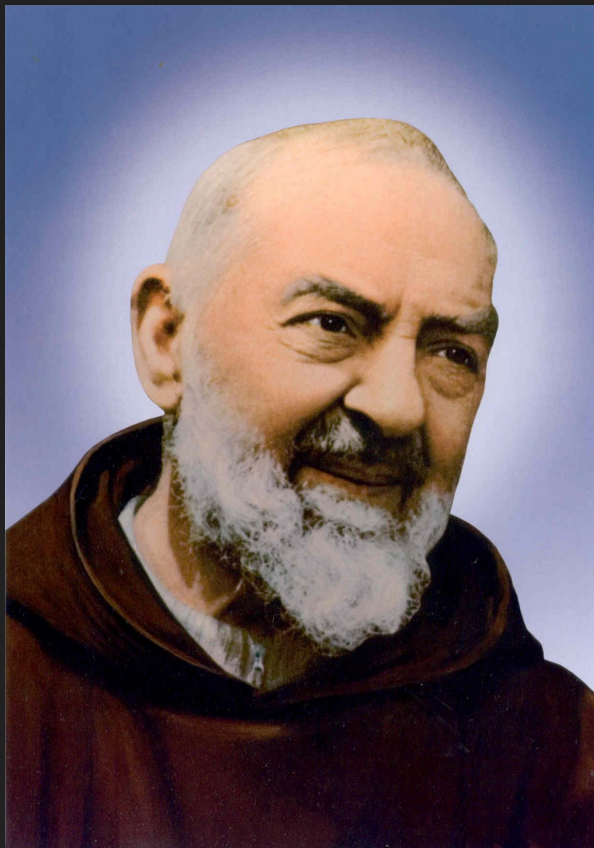
See

<https://mudongliang.github.io/2016/07/11/relo-a-not-so-well-known-memory-corruption-mitigation-technique.html>

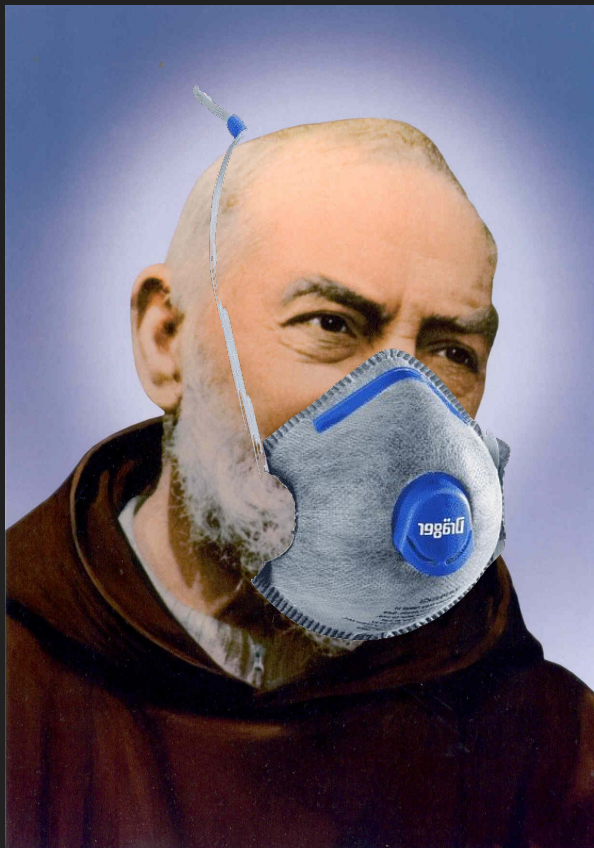
- Control flow integrity

See <https://nebelwelt.net/blog/20160913-ControlFlowIntegrity.html>

# Ultimate Protection



# Ultimate Protection 2020 version





# Now it's your turn!



<https://cyberchallenge.diag.uniroma1.it>



```
pip2 install ropper
```

```
pip2 install pwntools
```

```
install gef: https://gef.readthedocs.io/en/master/
```